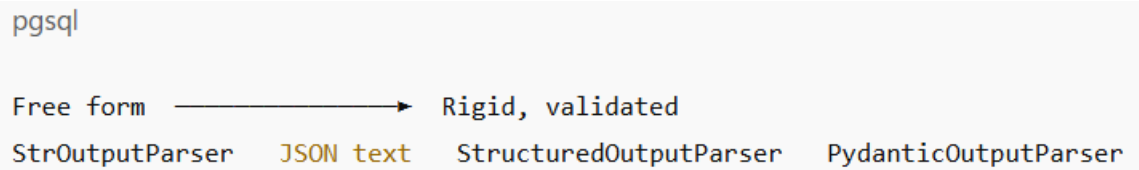


1) What Output Parsers are (and why they matter)

Output Parsers in LangChain turn raw LLM text into structured things your app can actually use—strings, dicts/JSON, or typed objects. They help with:

- **Consistency** – you always get a predictable shape back.
- **Validation** – you can catch missing/wrong fields early.
- **Ease of use** – no custom regex or manual json.loads gymnastics.

Think of them as a spectrum:



2) StrOutputParser (plain text)

What it is: The simplest parser; it just returns the model's output as a string.

When to use: Quick prototypes, natural-language answers, or when you truly don't need structure.

Pros

- Zero setup.
- Great for summaries, explanations, etc.

Cons

- No structure, no validation. You do all parsing yourself if you need data.

Snippet

```
python

from langchain_openai import ChatOpenAI
from langchain_core.prompts import ChatPromptTemplate
from langchain.output_parsers import StrOutputParser

llm = ChatOpenAI(temperature=0)
prompt = ChatPromptTemplate.from_template("Explain transformers in 3 sentences.")

chain = prompt | llm | StrOutputParser()
text = chain.invoke({})
```

3) “JSON output format” (instructed JSON, but no enforced schema)

What it is: You tell the model in the prompt “respond as JSON,” then parse it yourself (or with `JsonOutputParser`). This improves consistency, **but it does not enforce a schema**—the model can still miss fields, add extras, or use wrong types.

When to use: You want JSON quickly, don’t need strict guarantees, and are okay handling occasional format hiccups.

Pros

- Easy, flexible, minimal setup.
- Works fine for low-stakes data.

Cons

- **No guarantees:** fields may be missing/renamed/typed wrong.
- You must add your own checks.

Snippet

```
python

import json
from langchain_openai import ChatOpenAI
from langchain_core.prompts import ChatPromptTemplate

llm = ChatOpenAI(temperature=0)

prompt = ChatPromptTemplate.from_template(
    "Extract title and three tags from the text.\n"
    "Return ONLY JSON with keys: title (str), tags (list[str]).\n\n"
    "Text: {text}"
)

resp = (prompt | llm).invoke({"text": "LangChain parsers keep outputs tidy."})
data = json.loads(resp.content) # May throw; may not match schema perfectly
```

4) StructuredOutputParser (schema-shaped JSON via instructions)

What it is: A middle ground. You define **field names + human descriptions** via `ResponseSchema`. The parser:

1. Produces precise “format instructions” that you insert into the prompt.
2. Parses the result into a Python dict.

Validation level: Basic. It **nudges** the model to follow your schema (right keys), but types are still “stringly”; it won’t strongly validate numbers, enums, or nested types. Good for getting reliably-keyed JSON.

When to use: You want predictable keys and shape, but don’t need full Pydantic validation.

Pros

- Generates clear format instructions for the model.
- Good balance of ease and structure.

Cons

- Limited validation (presence/shape, not deep type checks).
- Complex/nested schemas are possible, but less ergonomic than Pydantic.

Snippet

```
python

from langchain_openai import ChatOpenAI
from langchain_core.prompts import ChatPromptTemplate
from langchain.output_parsers import StructuredOutputParser, ResponseSchema

schemas = [
    ResponseSchema(name="title", description="short headline for the text"),
    ResponseSchema(name="tags", description="list of exactly 3 short tags"),
    ResponseSchema(name="confidence", description="float between 0 and 1"),
]

parser = StructuredOutputParser.from_response_schemas(schemas)
format_instructions = parser.get_format_instructions()

prompt = ChatPromptTemplate.from_template(
    "Summarize the topic.\n{format_instructions}\n\nTopic: {topic}"
)

llm = ChatOpenAI(temperature=0)
chain = prompt | llm | parser

result = chain.invoke({"topic": "Output parsers in LangChain",
                       "format_instructions": format_instructions})
```

5) PydanticOutputParser (strict, type-safe, validated)

What it is: The most robust option. You define a **Pydantic model** (fields, types, defaults, validators). The parser:

1. Emits format instructions for the LLM based on your model.
2. Parses and **validates** the LLM output into a real Pydantic object (or raises `ValidationError`).

When to use: Production workflows, integrations, or any place you need **schema enforcement, type safety, and clean Python objects**.

Pros

- **Strict schema enforcement** (missing/extra/wrong types → errors).
- Built-in type conversion and custom validators (ranges, enums, nested objects).
- Plays nicely with the rest of your Python code.

Cons

- Slightly more setup.
- You should handle validation errors (and optionally auto-repair).

Snippet (with validation)

```
from typing import List
from langchain_openai import ChatOpenAI
from langchain_core.prompts import ChatPromptTemplate
from langchain.output_parsers import PydanticOutputParser
from langchain.pydantic_v1 import BaseModel, Field, ValidationError, validator

class Product(BaseModel):
    title: str = Field(..., description="Short product name")
    price: float = Field(..., description="Price in USD")
    tags: List[str] = Field(default_factory=list, description="3 short tags")
    rating: int = Field(..., description="1 to 5")

    @validator("rating")
    def rating_range(cls, v):
        if not (1 <= v <= 5):
            raise ValueError("rating must be 1..5")
        return v

parser = PydanticOutputParser(pydantic_object=Product)
format_instructions = parser.get_format_instructions()

prompt = ChatPromptTemplate.from_template(
    "Extract product info from the text.\n{format_instructions}\n\nText: {text}"
)

llm = ChatOpenAI(temperature=0)
chain = prompt | llm | parser

try:
    product: Product = chain.invoke({
        "text": "Name: Spark Widget; Price: 19.99; Tags: tools, diy, gadgets; Rating: 5",
        "format_instructions": format_instructions
    })
except ValidationError as e:
    # Handle or auto-fix (see next section)
    raise
```

6) Repairing/Hardening (useful add-ons)

- **OutputFixingParser:** If the model returns something slightly off (e.g., a trailing comma), this wrapper asks the LLM to “repair” it to satisfy your parser.
- **RetryOutputParser:** Automatically re-prompts the LLM with the validation error to try again.

Snippet (auto-fix with Pydantic)

```
from langchain.output_parsers import OutputFixingParser

fixing_parser = OutputFixingParser.from_llm(parser=parser, llm=llm)
chain = prompt | llm | fixing_parser
product = chain.invoke({
    "text": "...",
    "format_instructions": format_instructions
})
```

7) Choosing the right tool (quick guide)

- **Need plain text?** → StrOutputParser.
- **Want JSON but okay if it's imperfect?** → Instruct “return JSON” + json.loads or JsonOutputParser.
- **Want stable keys & shape, light checks?** → StructuredOutputParser.
- **Need real validation & Python objects?** → PydanticOutputParser (best for production).

8) Best practices (applies to all structured options)

1. **Always insert the parser's get_format_instructions() into your prompt.** That's how the LLM learns the expected shape.
2. **Keep temperature low (0–0.3)** when you need strict structure.
3. **Few-shot examples help.** Show a correct input → output pair, especially for nested schemas.
4. **Guard rails:** Wrap with OutputFixingParser or RetryOutputParser for resilience.
5. **Be explicit with types and ranges.** (e.g., “float 0–1”, “exactly 3 tags”).
6. **Handle errors.** Catch ValidationError and decide: retry, fix, or surface to the user.
7. **Prefer Pydantic for nested/typed data** (enums, unions, dates, custom validation).